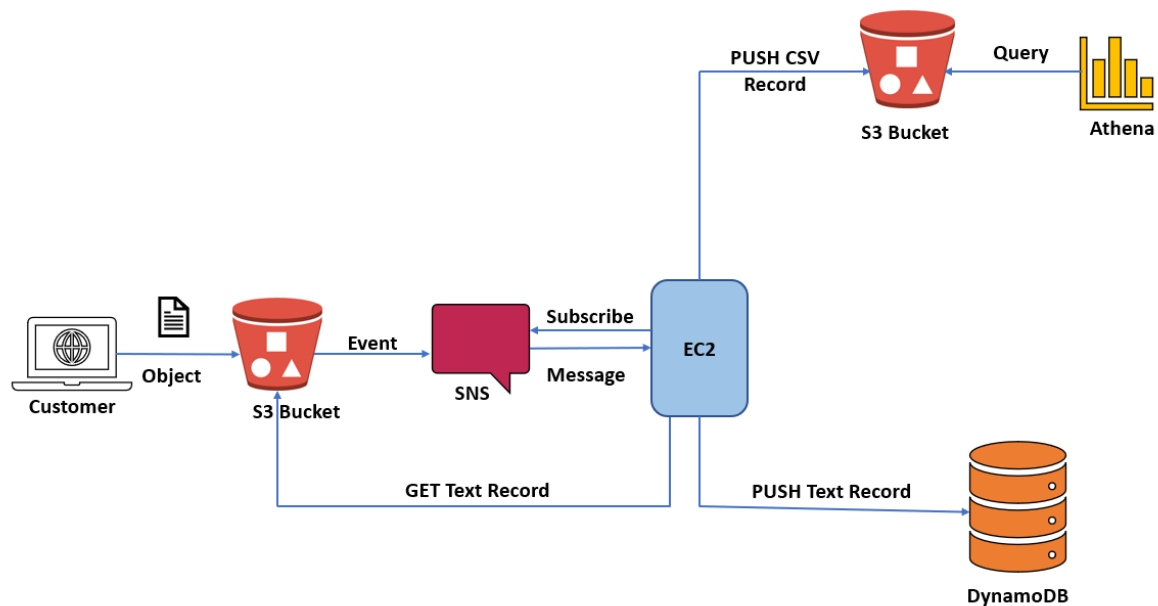


Creating an event triggered business process leveraging multiple managed services from AWS (EC2, S3, SNS, Athena, DynamoDB)

The objective of this project is to create an automated, event based real time process. Data should flow rapidly from the source to the destination. We will use and leverage multiple managed services available from a public cloud platform (AWS in this case) to achieve this.

ARCHITECTURE



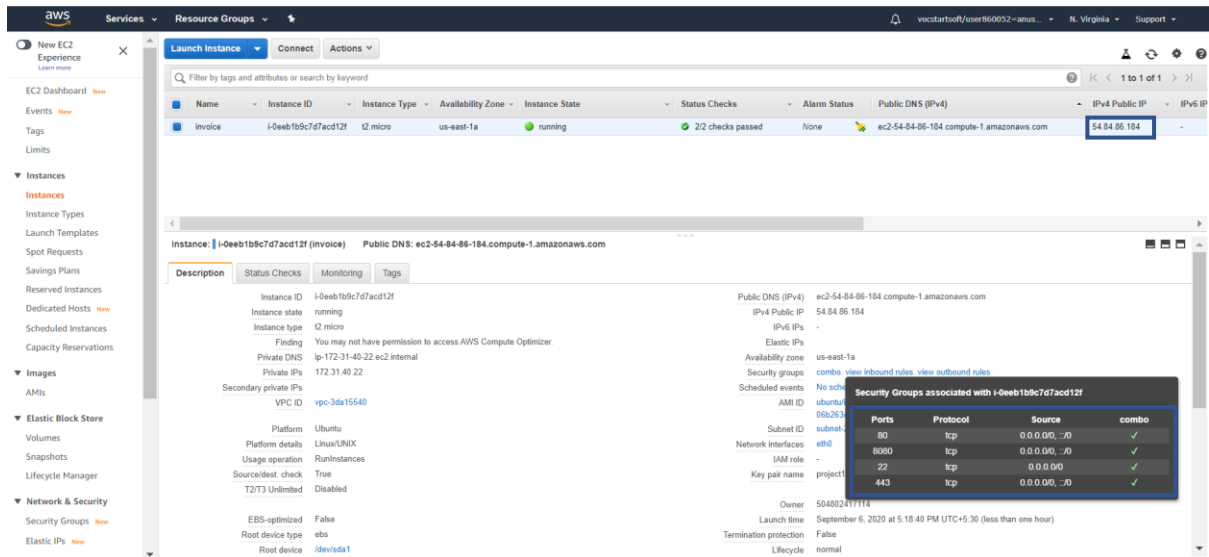
IMPLEMENTATION

The business process will be developed completely on the AWS public cloud using multiple managed services with some custom code for data transformation.

1. The customer uploads the invoice data to S3 bucket in a text format as per the organizations guidelines and policies. This bucket will have a policy to auto delete any content that is more than 1 day old (24 hours).
2. An event will trigger in the bucket that will place a message in SNS topic.
3. A custom program running in EC2 will subscribe to the SNS topic and get the message placed by S3 event.
4. The program will use S3 API to read from the bucket, parse the content of the file and create a CSV record along with saving the original record in DynamoDB.
5. The program will use S3 API to write CSV record to destination S3 bucket as a new S3 object.
6. Athena is used to query the CSV file (query to show aggregated expenses grouped by date).

IMPLEMENTATION

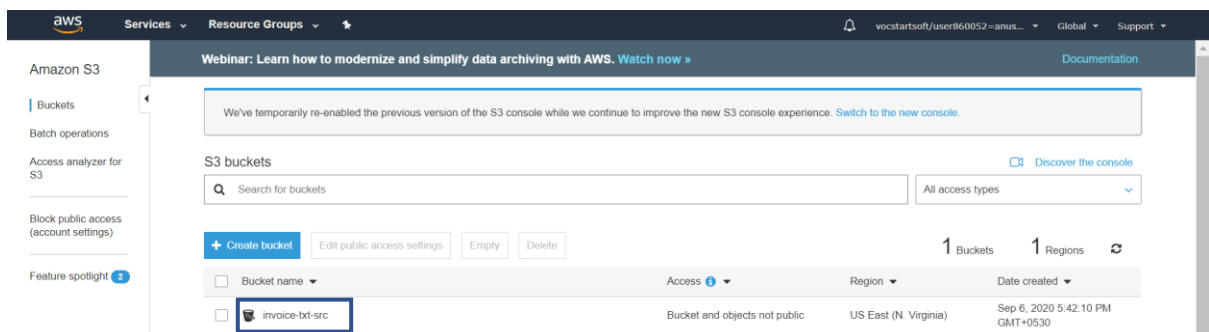
- 1) Create an Ubuntu EC2 instance and copy the custom code used for data transformation from a TEXT record to a CSV record

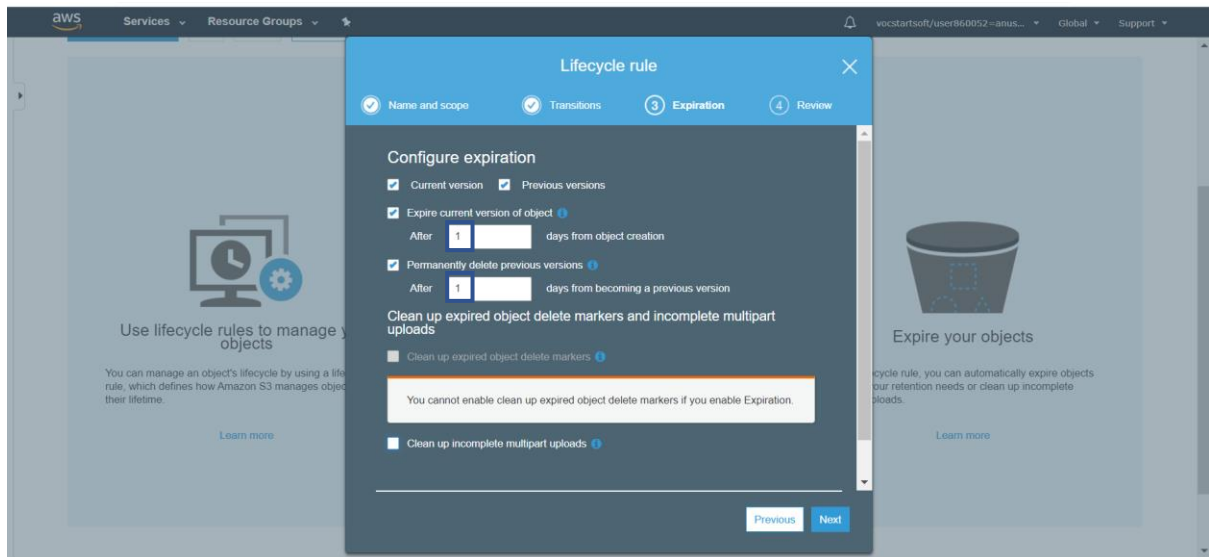


S3 Destination Bucket Name: invoice-csv-dest
DynamoDB Region: us-east-1


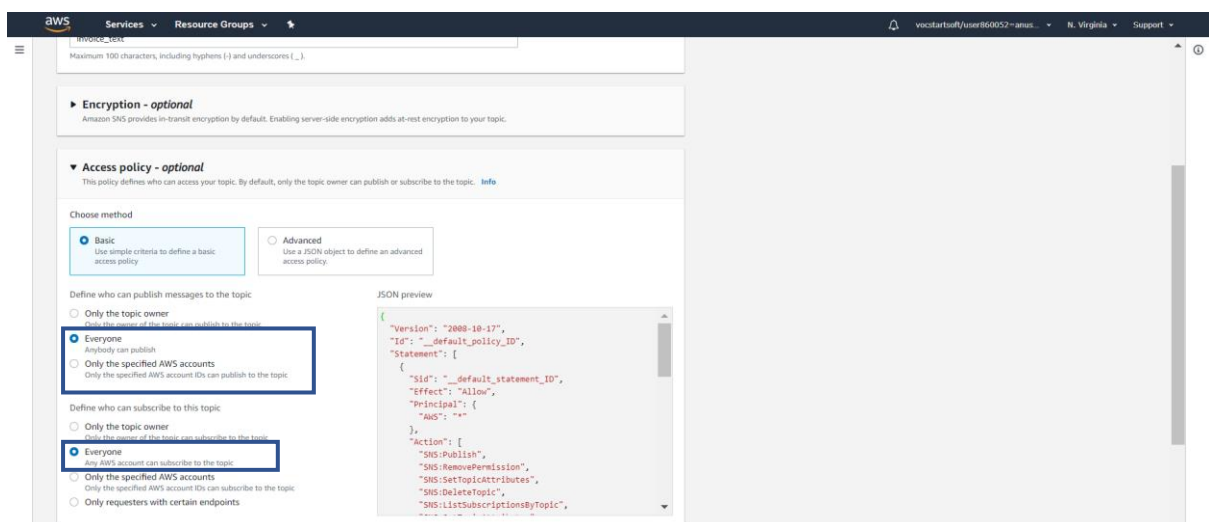
```
(djangoev) ubuntu@ip-172-31-40-22:/opt/docproc-new/api$ cat views.py | egrep -w "us-east-1|invoice-csv-dest"
dynamodb = boto3.resource('dynamodb', region_name='us-east-1')
s3_target_bucket = 'invoice-csv-dest'
(djangoev) ubuntu@ip-172-31-40-22:/opt/docproc-new/api$
```

- 2) Create a source bucket in S3 and add a rule to auto delete any object that is more than 1 day old (24 hours).
S3 Source Bucket: invoice-txt-src

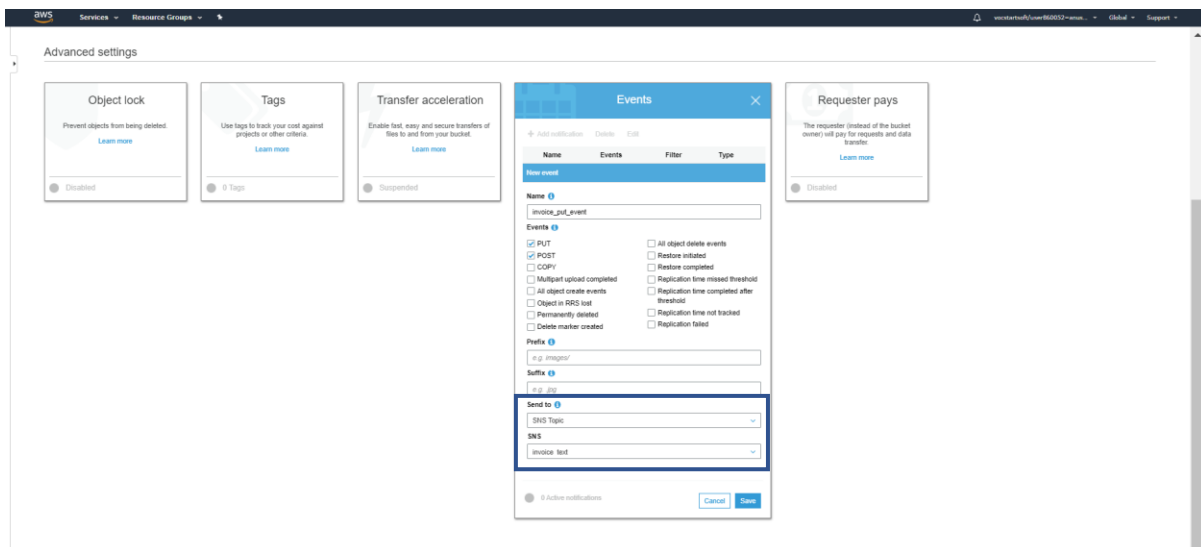




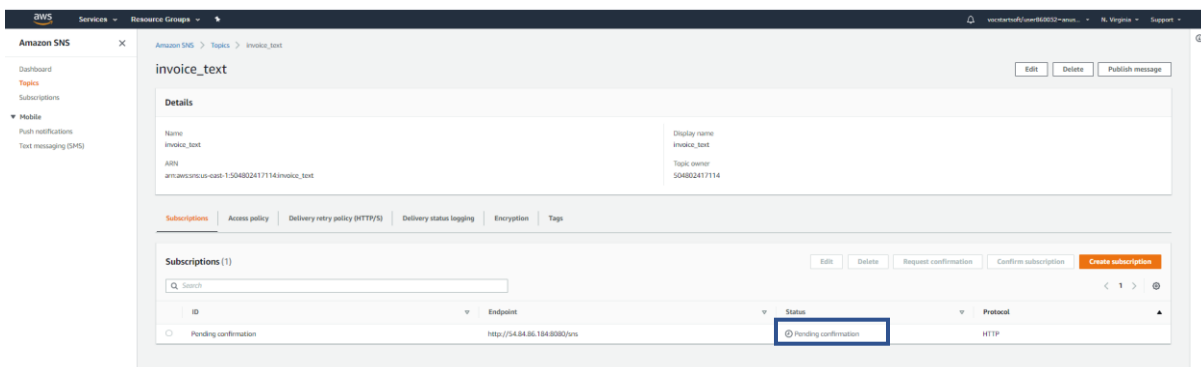
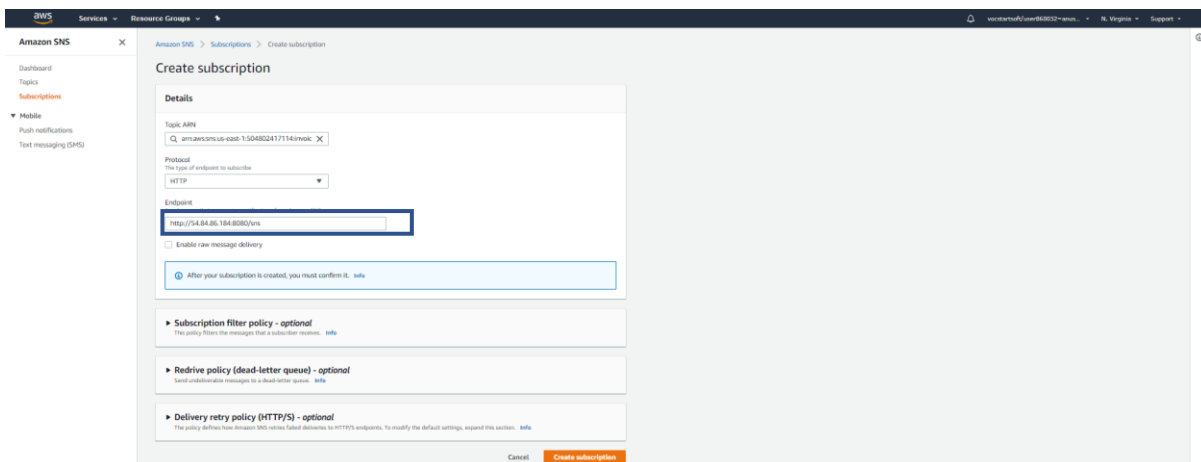
3) Create SNS topic Topic: invoice_text



4) Create an event in S3 that will place a message in SNS topic Event: invoice_put_event



5) EC2 application must subscribe for the SNS topic

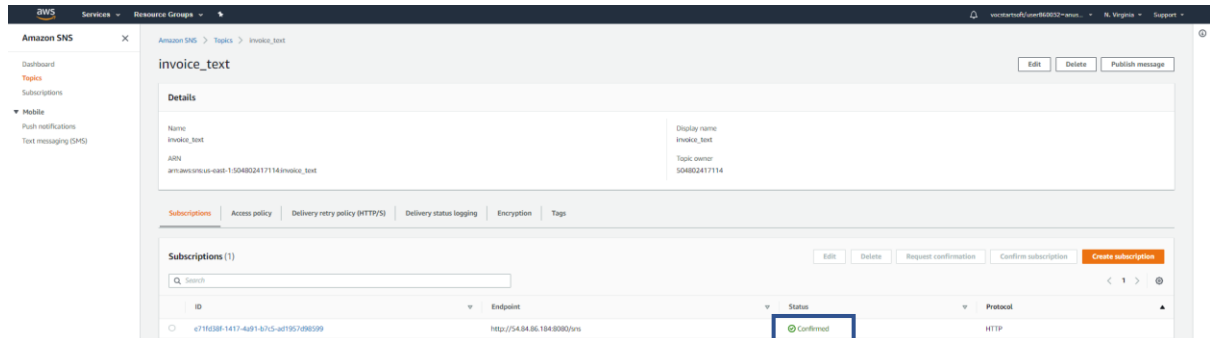


```
Request method is POST
The S3 JSON is {
  "Type": "SubscriptionConfirmation",
  "MessageId": "8053eb4d-8c37-4ef1-a806-cc083676f5df",
  "Token": "2336412f37fb687f5d51e6e2425f004ae03ffbe2e3ac3afc1fc8c327c70e2a8e8300bc5860fa8dda3b1c4fd2b39d2841a9daff2d6241c849b57725907cdfcfcc1659fea57eef1dd8d63f71ab47c3afa15220ec4e92289f62166ca96b411602afa3b889234981001a0dd7120186cb37",
  "TopicArn": "arn:aws:sns:us-east-1:504802417114:invoice_text",
  "Message": "You have chosen to subscribe to the topic arn:aws:sns:us-east-1:504802417114:invoice_text.\nTo confirm the subscription, visit the SubscribeURL included in this message.",
  "SubscribeURL": "https://sns.us-east-1.amazonaws.com/?Action=ConfirmSubscription&TopicArn=arn:aws:sns:us-east-1:504802417114:invoice_text&Token=2336412f37fb687f5d51e6e2425f004ae03ffbe2e3ac3afc1fc8c327c70e2a8e8300bc5860fa8dda3b1c4fd2b39d2841a9daff2d6241c849b57725907cdfcfcc1659fea57eef1dd8d63f71ab47c3afa15220ec4e92289f62166ca96b411602afa3b889234981001a0dd7120186cb37",
  "Timestamp": "2020-09-06T12:35:13.679Z",
  "SignatureVersion": "1",
  "Signature": "YPa7TuYEJgZJaZCeGZAmX6Q9qnLj67miSrZHDQDxKE33ZohxF/Yb6eqp6Lb95SuV9e4NoGjuIMS/6UJwc5faLaTpETapH/SjikAJBm6LZ6eb0m2dR5/W EymHyE6U6iQ+zhV9q/KSB80Ds0iZNIk9cp8UxoBHxqSKHdwXVBTKLGmp/NmPZUzrT86u9s51u0dR+3ievIL3wn3rQLvP/RwUSx9Lms6xhsPPU2GU6oreQBjx0EgPVii6+BA+UqoEcjntXCdwVp+j06B8jZb6LbthniN/oAQIy6uQiyvIYHYwVbWHC+nCw0FRAQLDcZm4rQoToIHmyF1IJO5mfr9LURCg==",
  "SigningCertURL": "https://sns.us-east-1.amazonaws.com/SimpleNotificationService-a86cb10b4e1f29c941702d737128f7b6.pem"
}
```

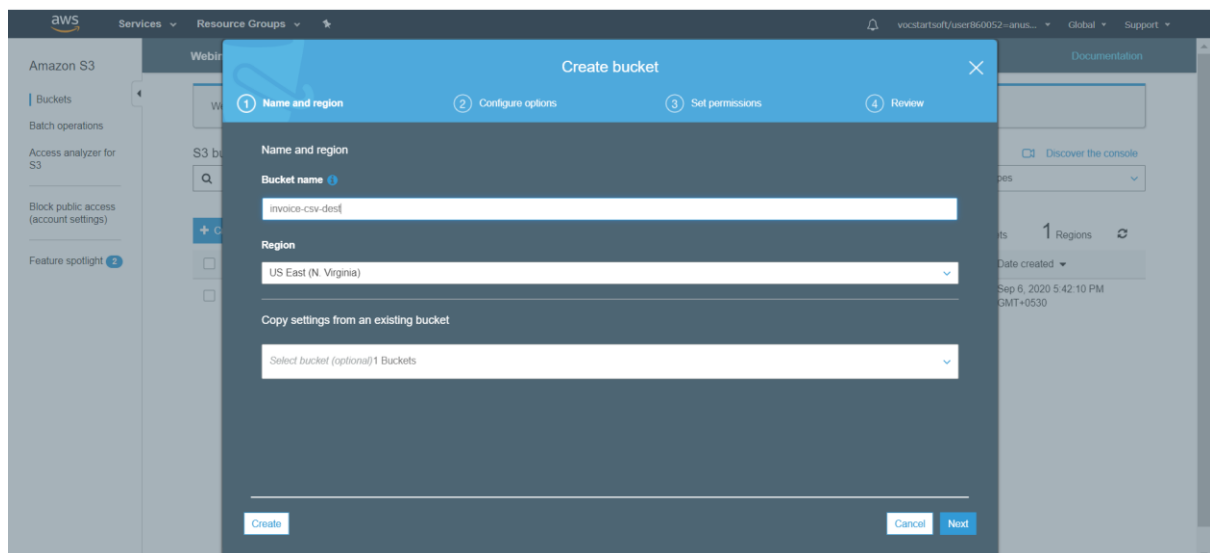
← → ↻ 🔒 sns.us-east-1.amazonaws.com/?Action=ConfirmSubscription&TopicArn=arn:aws:sns:us-east-1:504802417114:invoice_text&Token=2336412f37fb687f5d51... ☆ 🛠️ 👤 ⋮

This XML file does not appear to have any style information associated with it. The document tree is shown below.

```
<?xml version="1.0" encoding="UTF-8"?>
<ConfirmSubscriptionResponse xmlns="http://sns.amazonaws.com/doc/2010-03-31/">
  <ConfirmSubscriptionResult>
    <SubscriptionArn>arn:aws:sns:us-east-1:504802417114:invoice_text:e71fd38f-1417-4a91-b7c5-ad1957d98599</SubscriptionArn>
  </ConfirmSubscriptionResult>
  <ResponseMetadata>
    <RequestId>fc1e5af2-7f39-5a7c-b72d-8d3ca5f94db4</RequestId>
    </ResponseMetadata>
  </ConfirmSubscriptionResponse>
```



6) Create destination bucket in S3 which will hold the CSV records



7) Create an IAM role to add S3FullAccess and DynamoDBFullAccess, attach the role to the EC2 instance

aws Services Resource Groups

Create role

Review

Provide the required information below and review this role before you create it.

Role name: invoice-role

Role description: Allows EC2 instances to call AWS services on your behalf.

Trusted entities: AWS service: ec2.amazonaws.com

Policies: AmazonS3FullAccess, AmazonDynamoDBFullAccess

Permissions boundary: Permissions boundary is not set

The new role will receive the following tag

Key	Value
Name	invoice

Cancel Previous Create role

aws Services Resource Groups

Instances > Attach/Replace IAM Role

Attach/Replace IAM Role

Select an IAM role to attach to your instance. If you don't have any IAM roles, choose Create new IAM role to create a role in the IAM console. If an IAM role is already attached to your instance, the IAM role you choose will replace the existing role.

Instance ID: i-0eeb1b9c7d7acd12f (invoice)

IAM role: invoice-role

Create new IAM role

Cancel Apply

8) Upload a TEXT invoice document to S3 source bucket

docproc-invoice - Notepad

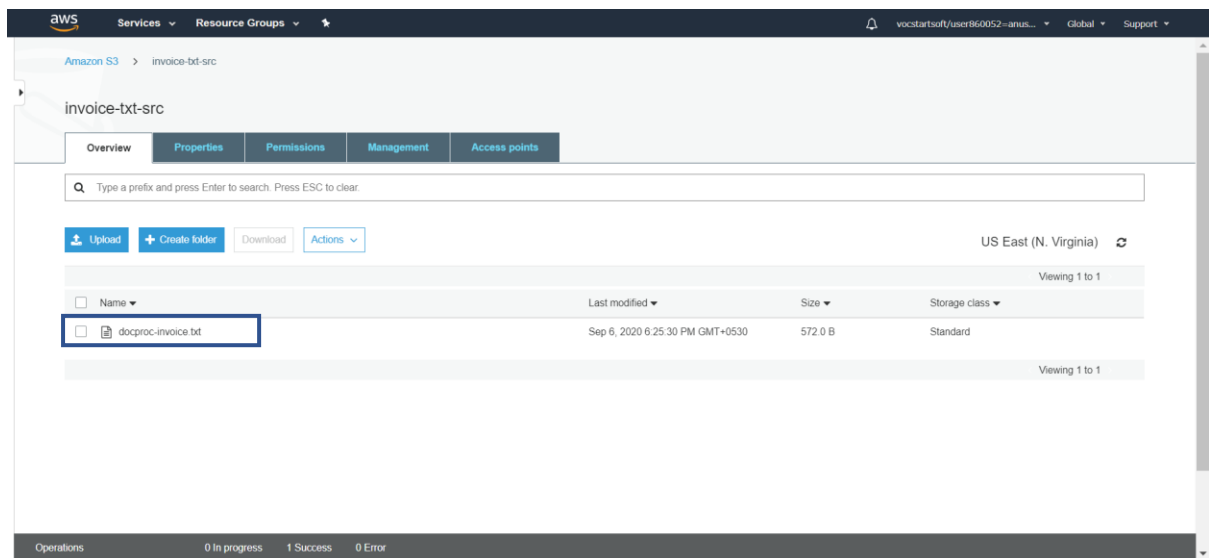
File Edit Format View Help

Invoice

Customer-ID: sm-LK0080145
 Inv-ID: inv-00001
 Dated: Mar 31 2018
 From: Space Mechanics Inc.
 To: Starfleet Interplanetary HQ
 Amount: 7650
 SGST: 500
 Total: 8150
 InWords: Space Rupees Eight thousand one hundred fifty only

Items

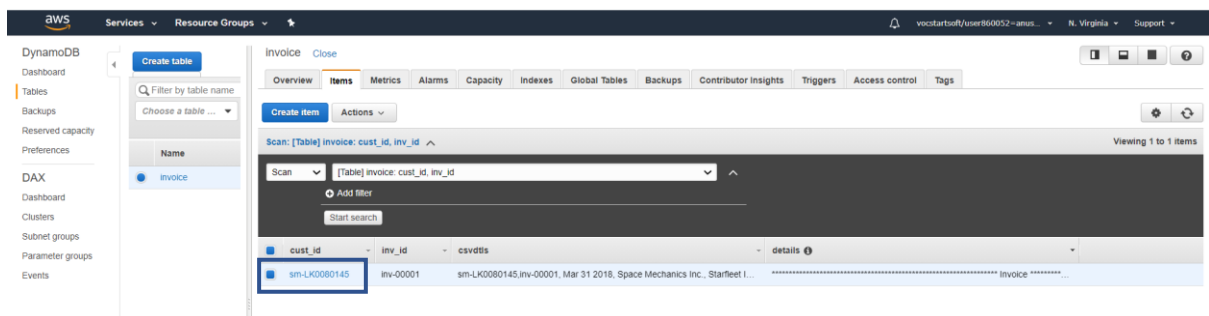
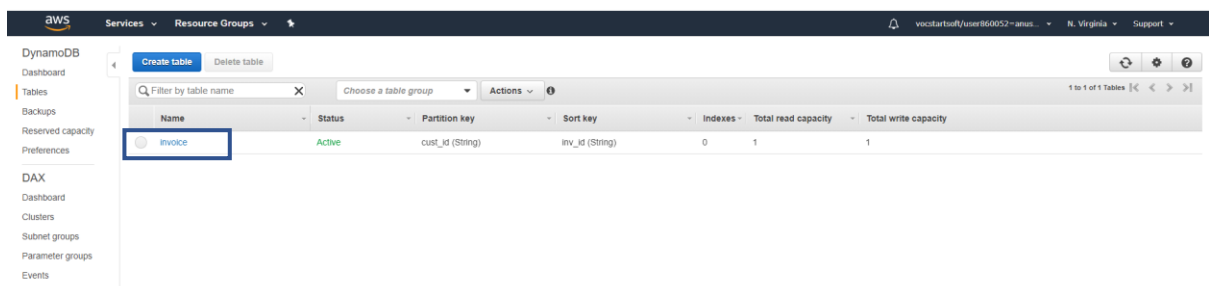
- 1) Matter antimatter fusion controller, 2000, 200
- 2) External inertial damper, 3100, 180
- 3) Type X phaser, 500, 35
- 4) 2MW Fusion reactor, 1200, 85
- 5) Transporter base dial, 850



9) Navigate to DynamoDB and see the TEXT invoice record placed

A New table named invoice is created and a record is placed in the table

```
*****
Creating table invoice
Skipped due to exception ResourceInUseException
Reason Table already exists: invoice
*****
Inserting data in the table
Done.
Written new s3 file 10297710_sm-LK0080145_inv-00001.csv
Uploading S3 object content sm-LK0080145_inv-00001, Mar 31 2016, Space Mechanics Inc., Starfleet Interplanetary HQ, 7650, 500, 8150,
Space Rupees Eight thousand one hundred fifty only
Done
[06/Sep/2020 12:58:10] "POST /sns HTTP/1.1" 200 43
DONE
*****
```

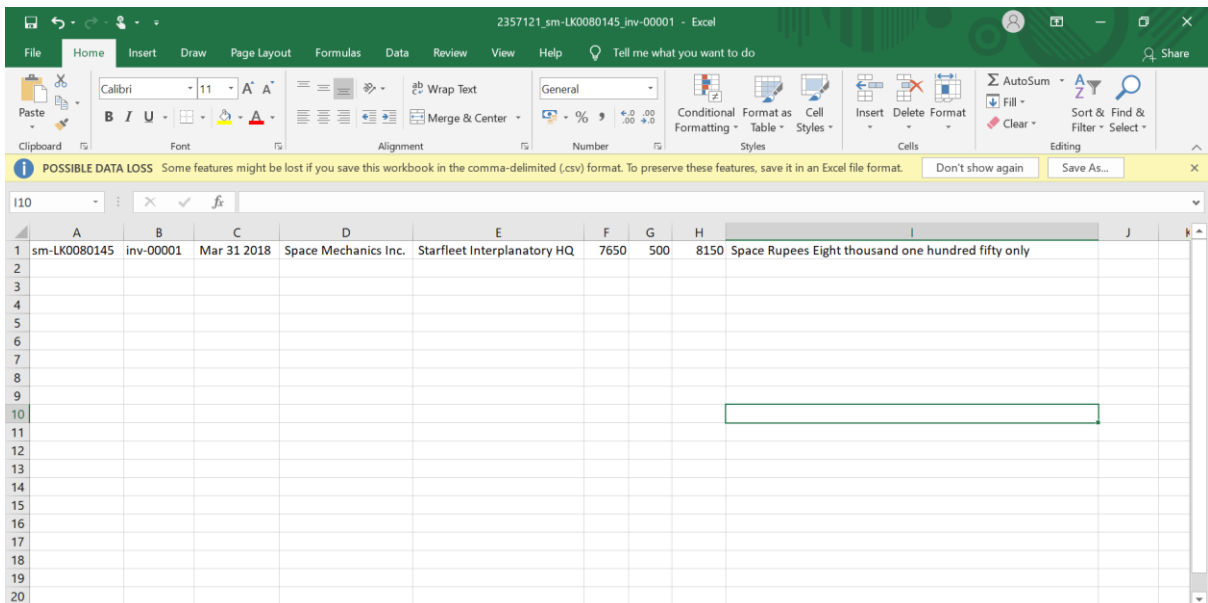
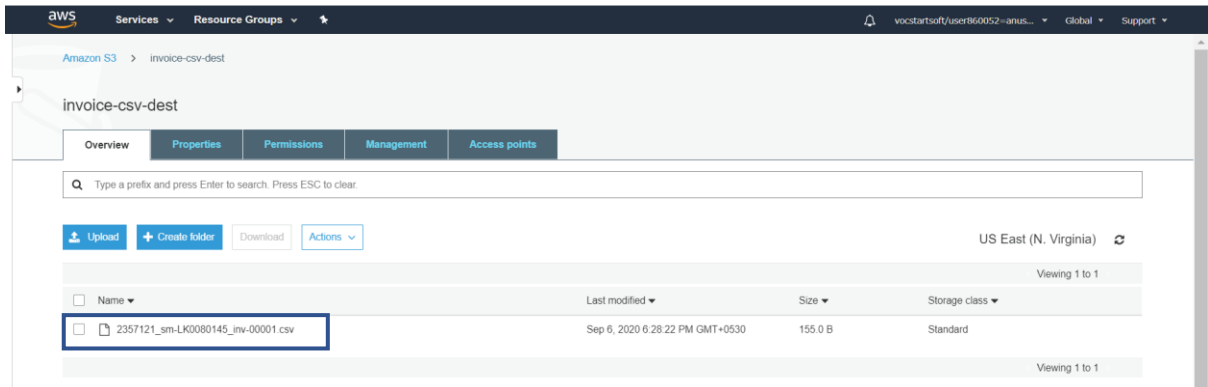


10) See the CSV record created in S3 Destination bucket

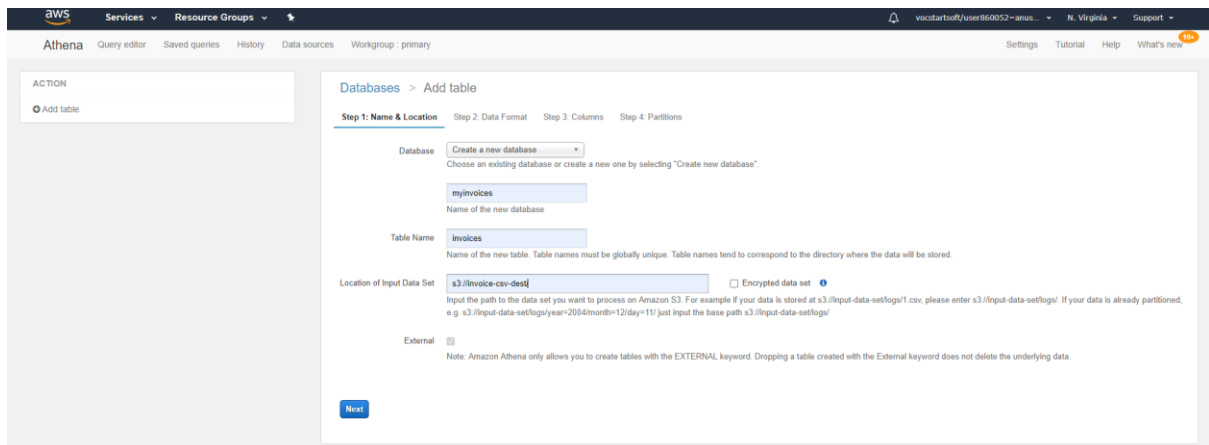
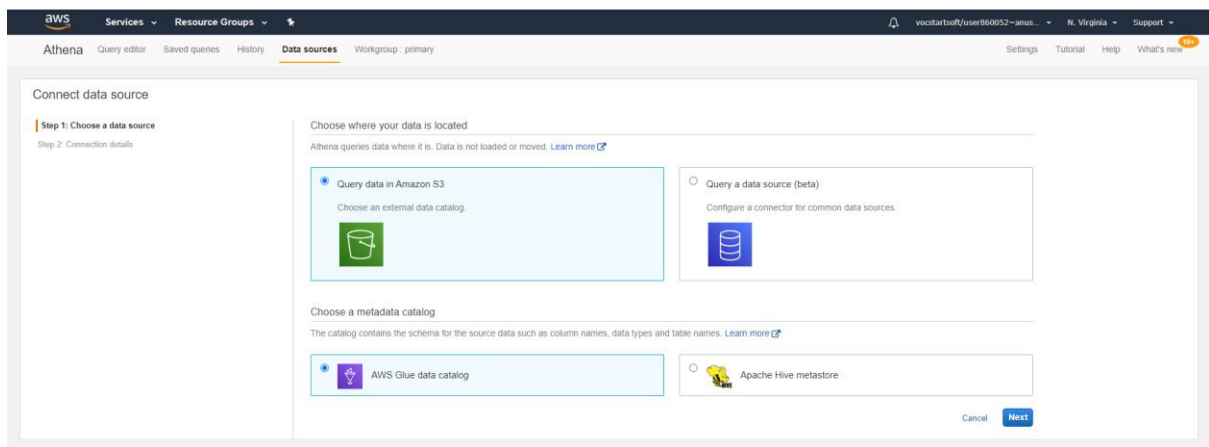
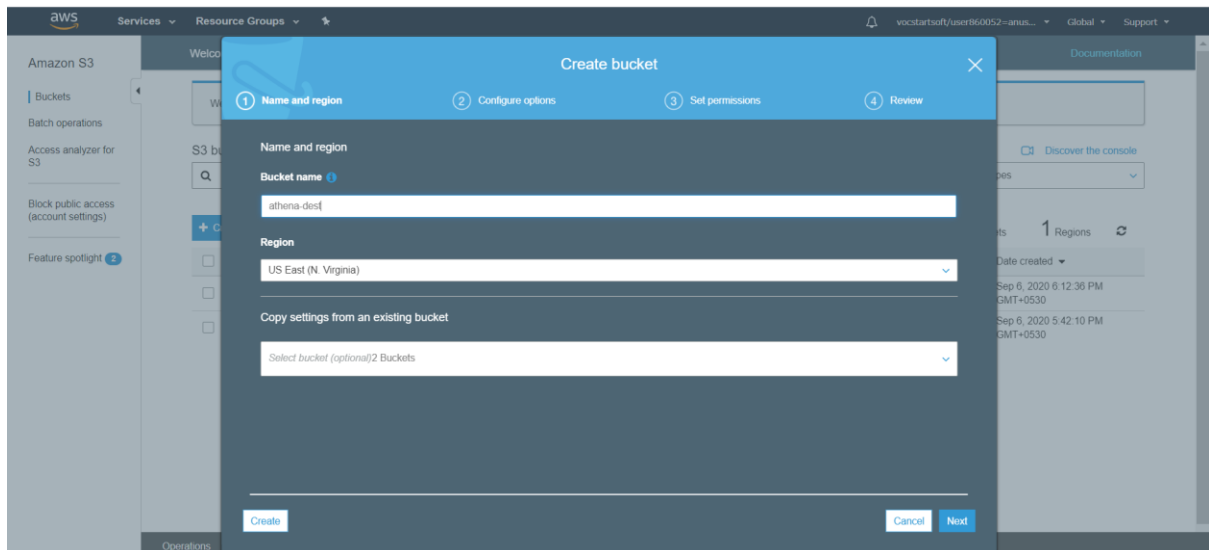
```

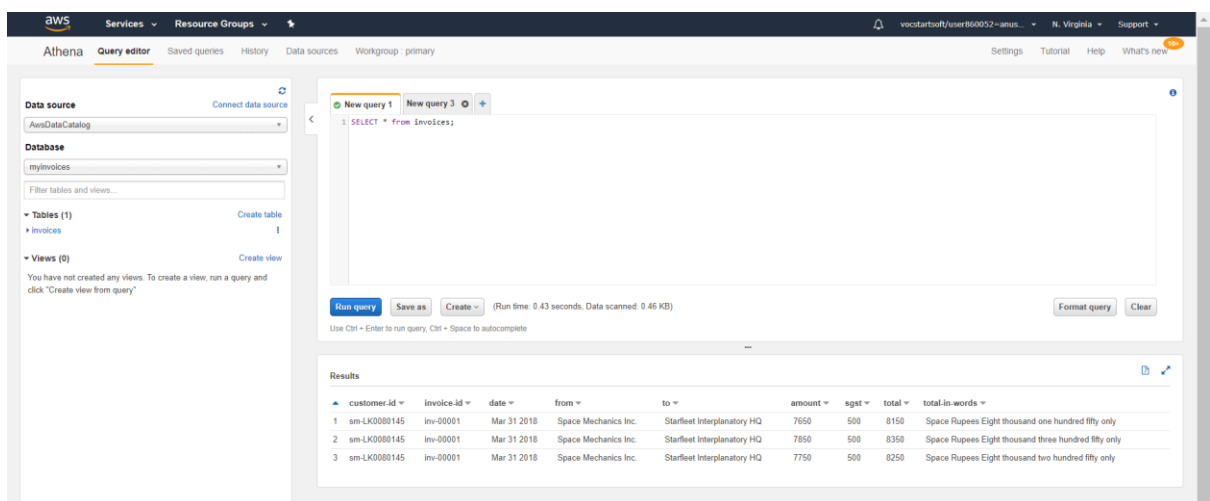
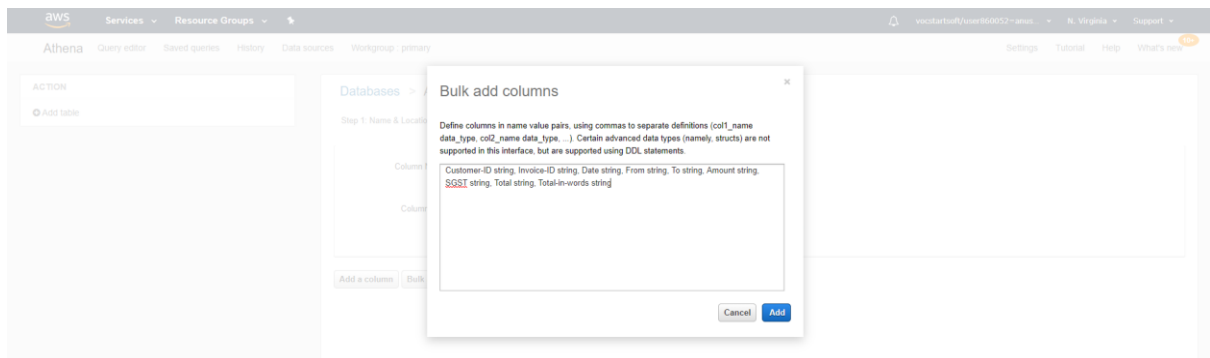
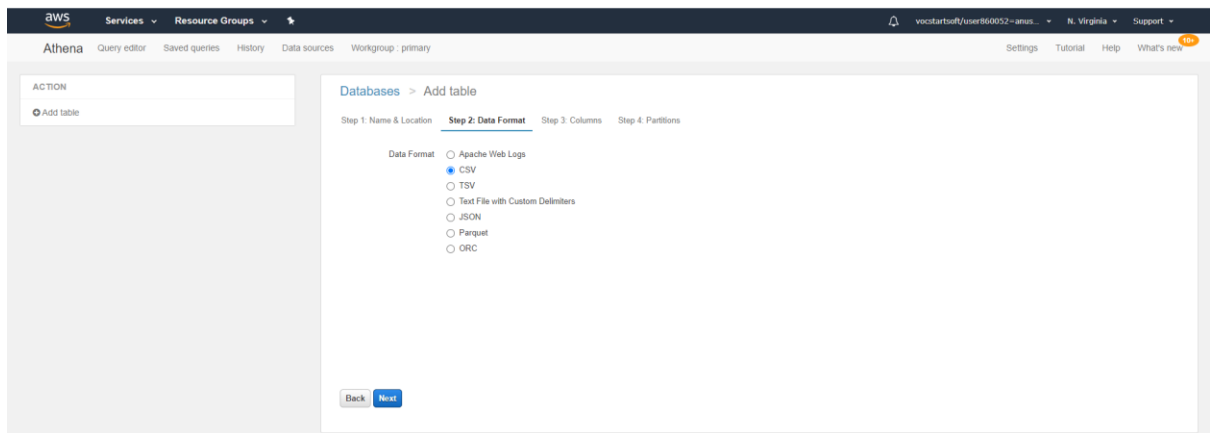
Will read the file docproc-invoice.txt from the bucket invoice-txt-src
Line-> *****
Line-> Invoice
Line-> *****
Line-> Customer-ID: sm-LK0080145
Found Customer-ID sm-LK0080145
Line-> Inv-ID: inv-00001
Found Invoice-ID inv-00001
Line-> Dated: Mar 31 2018
Line-> From: Space Mechanics Inc.
Line-> To: Starfleet Interplanetary HQ
Line-> Amount: 7650
Line-> SGST: 500
Line-> Total: 8150
Line-> InWords: Space Rupees Eight thousand one hundred fifty only
Line->
Line-> Items
Line-> 1) Matter antimatter fusion controller; 2000; 200
Line-> 2) External inertial damper; 3100; 180
Line-> 3) Type X phaser; 500; 35
Line-> 4) 2MW Fusion reactor; 1200; 85
Line-> 5) Transporter base dial; 850
Found dated Mar 31 2018
Found fromcust Space Mechanics Inc.
Found tocust Starfleet Interplanetary HQ
Found amt 7650
Found sgst 500
Found tot 8150
Found words Space Rupees Eight thousand one hundred fifty only
CSV -> sm-LK0080145,inv-00001, Mar 31 2018, Space Mechanics Inc., Starfleet Interplanetary HQ, 7650, 500, 8150, Space Rupees Eight thousand one hundred fifty only

```



11) Configure Athena to query the CSV record in S3 destination bucket





aws

Services

Resource Groups

Account ID: wccstartsoft/user66052-anus

Region: N. Virginia

Support

Athena

Query editor

Saved queries

History

Data sources

Workgroup: primary

Settings

Tutorial

Help

What's new

Data source

Connect data source

AwsDataCatalog

Database

myinvoices

Filter tables and views...

Tables (1)

Create table

Invoices

Views (0)

Create view

You have not created any views. To create a view, run a query and click "Create view from query"

New query 1

New query 3

+

1

SELECT date, MAX(total) FROM Invoices GROUP BY date;

Run query

Save as

Create

(Run time: 0.5 seconds, Data scanned: 0.46 KB)

Format query

Clear

Use Ctrl + Enter to run query, Ctrl + Space to autocomplete

Results

date

_col1

1

Mar 31 2018

8350